

Mastery Learning in Practice: A Descriptive Analysis of Data from the Cognitive Tutor Algebra I Effectiveness Trial

Anita Israni

Adam C Sales

John F Pane

October 23, 2017

1 Introduction

Mastery learning sits at the foundation of intelligent tutoring systems. The philosophy of mastery learning assumes a well-structured curriculum, and posits that students progress within the curriculum as they master its skills. The Cognitive Tutor Algebra I (CTAI) system, developed by Carnegie Learning, Inc., is one of the best studied and best regarded examples of modern educational software. It is a blended learning system for teaching algebraic concepts and principles, to middle and high school students, including both text-book materials and software. The software component of the curriculum allows students to progress at their own pace and receive individualized feedback on their performance. A large-scale randomized effectiveness trial conducted by the RAND corporation showed that, in some circumstances, CTAI boosts students scores on an Algebra-I post-test by about one fifth of a standard deviation (Pane, et al. 2013). CTAs success in this experiment would seem to vindicate its pedagogy: mastery learning, and the algebra I curriculum on which it is based.

However, the theory underlying CTAI does not always determine its use. To be sure, the software has a standard set of algebra topics, divided into units and further into sections; and a standard sequence for presenting them. But this precise curriculum is not mandatory. At the request of educators, it can be customized by altering what units or sections are included (including, possibly, material from a different standard curriculum such as geometry), as well as their sequence, to conform to local or state standards. Further, teachers have the option of moving students within the curriculum, regardless of the softwares estimate of their skill mastery.

This article examines teachers and schools ^{adherence or} non-adherence to the standard, mastery-based CTAI curriculum, using data from the RAND study, a seven-state randomized controlled trial of CTAI in high schools and middle schools. The study found a significant positive effect of CTAI in high schools, during their second year of implementation but not the first. Students in the treatment group of the study were enrolled in one or more of the

standard or customized ^{curricula} curriculum during their participation in the study, and the software logged aspects of their usage, including time spent, sections encountered, and whether the software judged the students to have mastered the sections. ^{Adopting standard procedures of an effectiveness trial, both} Both Carnegie Learning and the researchers running the study restricted their ~~software~~ support and oversight to what is typically provided outside of an experimental context. Thus, the data from the study reflect ^{software} typical ~~software~~ usage. Secondary data analyses used principal stratification to show that students who attempted more sections experienced larger treatment effects, and students who had high or low assistance levels, as opposed to an average level, experienced smaller treatment effects (Sales & Pane, 2015, Sales et al. 2016).

^{Not sure we should call this an extension of the prior work.} Here we ~~extend the previous findings to~~ investigate patterns of change in the sequence of topics presented to students. As noted earlier, schools could use a standard curriculum, a customized curriculum, or switch from one curriculum to another. Each curriculum consisted of a sequence of units, which was comprised of a sequence of sections. For example, the standard CTAI Algebra I curriculum in Year 1 comprised of 40 units covering approximately 170 sections. The number of sections within each unit varied. This study attempts to elucidate teachers' goals in reassigning students. One hypothesis is the need for teachers ^{to} push ahead students who were falling behind, i.e. to reassign them from sections on which they were struggling, to allow them catch up with the rest of the class. Another hypothesis is that teachers sought to push students past easier sections to begin working on more relevant ^{or challenging} topics for them. A third hypothesis is that teachers needed to cover certain topics in preparation for an upcoming state exam, and might have reassigned groups of students all at the same time to cover topics that might otherwise not have been covered. This hypothesis may lead to an increase in reassignments as the exam approaches. Overall, understanding patterns of reassignment is the focus of this paper, along with insights into how teachers were using the software component of CTAI.

2 The RAND Effectiveness Trial

The study to measure the effectiveness of CTAI included 7 states, 73 high schools, and 74 middle schools with nearly 18,700 high school students and 6,800 middle school students. ^{participating} Schools were enrolled in a total of 52 school districts that were distributed ^{among} between urban, suburban, and rural areas. Schools were matched on a set of covariates, and then randomly assigned to the treatment or control group. Schools in the control group continued with their current algebra curriculum, and schools in the treatment group used Carnegie Learning's curriculum which includes textbook materials as well as CTAI. Each school participated for two years, with a different cohort of students ^{participating} taking algebra the second year (with a small fraction of students present in the study both years because they repeated algebra). It should be noted that this study did not include statewide implementations; the study results cannot be generalized to all schools within the state. In some states, one large school district participated, while in other states, a set of smaller school districts participated. The states included Alabama (AL), Connecticut (CT), Kentucky (KY), Louisiana (LA), Michigan (MI), New Jersey (NJ), and Texas (TX). Each state participated in both the middle school and

↑ after having established ² these 2-letter state names, consider using them everywhere instead of spelling out the state names.




Figure 1: Percentage of worked problems coming from various levels (denoted by color, with Algebra II and Geometry bundled as “>Algebra I”), from standard and customized variants, denoted by shading.

high school arms of the study, except AL, which participated only in the middle school arm. The current study focuses on high school students only, due to the large number of students in this cohort and the significant positive treatment effect previously found for this cohort. All discussion that follows refers exclusively to the high school data.

There are some limitations to the available data for this study. In particular, the developer was able to provide the intended sequence of units/sections for only a subset of the customized curricula. Data from some schools, and some students within schools, were missing either because the log files were not retrievable, or because of an imperfect ability to link log data to other study data files. In addition, there were missing completion dates for many customized curricula; making it difficult to analyze the customized curricula when any sort of date was needed, such as time to complete a section.

3 Standard and Customized Curricula

Students’ automatic progress through the Cognitive Tutor (CT) software is governed by the predefined curriculum sequences of sections and units embedded in the software. Without external meddling, the curriculum a student works on determines what section he or she will be directed towards next after mastering (or exhausting the problems) from a previous section. In the CTAI effectiveness trial, the most common curriculum was, naturally, Algebra I. This came with three closely related variants, due to software releases. Students requiring more remediation were able to work on a less advanced curriculum, called “Bridge to Algebra,” and more advanced students could work on Algebra II or Geometry.

In the second year of the study, some high schools, primarily in Texas, Michigan, and Kentucky, requested customized variants of the curricula. This was typically due to state standards, ~~and~~ testing schedules. These “customized curricula” altered the order of some sections and units, and were ^{usually} particular to schools.

Figure 1 shows the percentage of worked problems from each curriculum, from standard and customized varieties, by state and year. First, note that the vast majority of worked problems were from the Algebra I sequence. A small but notable number of more advanced problems were worked in Kentucky in Year 2, and some less advanced problems were worked in Michigan and Louisiana. Secondly, note the rise in “customized curricula” in year 2 in Texas, Kentucky, and Michigan, the three states with the most students in our dataset. In particular, Texas shifted almost entirely to customized curricula from years 1 to 2.

Throughout the school year, teachers could have a class of students working on multiple curricula either sequentially, where the students changed curricula in lock step, or simultaneously, where students worked on different curricula at the same time. As an example, two teachers located in KY had their students working on Algebra I throughout most of the year and then reassigned them to Algebra II in the last month of school. Another teacher in KY had students variously enrolled in three different curricula throughout the entire year (Bridge-to-Algebra, Algebra I, and a specialized Geometry curriculum) while another teacher in MI during Year 2 had students enrolled in three curricula sequentially throughout the year covering Algebra I until November, and then starting a customized curriculum and using it until February, and ending the year with another customized curriculum that was used until June. While there are numerous instances of these uses of multiple curricula in Year 2, there are also many occurrences of teachers who had their students enrolled in the standard Algebra I throughout the entire year. In TX mostly, with some instances in MI, there were teachers who exclusively used customized curricula the whole year.

4 Student Usage Across States and Years

	Hours	Problems	Sections	Units
Year 1	33.47	309	43.0	9
Year 2	23.83	228	36.5	10

Table 1: Median numbers of hours, problems, sections, and units worked by each student in the dataset in the two years of the study. Students with no usage data were excluded.

Table ¹ ?? shows the median numbers of hours, problems, sections, and units worked on by each student in the dataset in the two years of the study. Apparently, usage decreased markedly in the second year: the median student worked about 10 fewer hours, 81 fewer problems, and 6 fewer sections in year 2 than ^{their counterpart} in year 1.

Figure 2 shows ~~that~~ the number of hours students spent working on the CT software in some more detail, via state-by-year boxplots. Analogous figures for the numbers of problems and sections students worked, available in an online supplement, showed similar patterns

Consider rewording “the median student” because most likely it was not one unique student who was at the median of all these metrics




Figure 2: Boxplots of hours each student spent on Cognitive Tutor software over by year and state. Students with no timestamp data (193), with negative times (3) or with more than 110 hours (91) were excluded.

for the numbers of sections and problems students worked. Usage time varied substantially between students and across states and years. Time spent varied substantially by state, with students in Texas, Connecticut, and New Jersey working far fewer hours than students in Kentucky, Louisiana, and Michigan. Not every state reduced its usage from years 1 to 2—while students in Texas, Kentucky and New Jersey used the software less in the second year than in the first, students in Michigan, Louisiana, and Connecticut increased their usage.

Overall, usage varied a bit more in year 2 than in year 1—the median absolute deviation of time spent was 15.9 hours in the first year, compared to 16.9 in the second year. The increase in variation seems to be driven both by increasing between-state variation, and a between-student increase in Louisiana. One intriguing possibility is that CT usage was better tailored to teachers and students in the second year than in the first—perhaps students who stood to gain more from the software increased their usage, and the rest of the students used it less.

In contrast to the numbers of hours, problems, and sections students worked, Table 1 shows that the median number of units students worked increased by 1 from year 1 to year 2. This suggests students in year 2 were exposed, on average, to a slightly wider range of topics. Figure 3 shows boxplots of the numbers of units worked by state and year. The geographic variation in units worked mirrors the pattern in Figure 2, with more usage in Kentucky, Michigan, and Louisiana but less in Texas, Connecticut, and New Jersey. However, in every state the median Year 2 student worked at least as many different units as the median Year 1 student. Variation in the number of units worked also increased slightly from years 1 to 2—the interquartile range increased in every state except for Kentucky, where a decrease in




Figure 3: Boxplots of the number of units of Cognitive Tutor software each student worked, by year and state. Students working more than 55 units (7 of 4460) and students with no usage data (189) were excluded.

IQR was accompanied by an increase in the number of outliers.

All in all, students used CT software less in year 2 than in year 1. On the other hand, students in the second year tended to see a slightly wider range of topics, and varied somewhat more in their usage.

5 Working Units in Order—Or Not

Overall, students used CT less in the second year than in the first. How was this difference distributed across between CTAI units?

Figure 4 shows the percentage of students with usage data who worked each section from the standard CTAI curriculum. The units are ordered, from left to right, as they are prescribed in the standard CTAI curriculum.

In year 1, the curve is almost monotonically decreasing, as one would expect if students adhered to the curriculum. Students varied in how many units they worked—with the variation due to both student ability and the amount of time allocated to CTAI within a classroom—but they mostly followed the same standard curriculum. Students who worked fewer units stopped earlier in the sequence, and those who worked more units progressed farther. Hence, earlier sections were worked by higher proportions of students than later units.

In contrast, in year 2, students were much more likely to buck the standard unit order. For instance, Figure 4 suggests that students skipped “Unit Conversions” and “1 step Lin. Eq.” for “1st Quadrant Linear Graphs” and “Ind. Variables in Lin. Mod.” Both of the latter sections were worked by a greater proportion of students than the former sections.

Figure 4 shows the units of algebra along the horizontal axis in the order they are prescribed. The vertical axis shows the percentage of students with usage data who worked each section.

depart from
suggested rewording on next page




Figure 4: The percentages of students with usage data who worked at least one problem from each unit in the Algebra I curriculum. The units are arranged in order for the standard curriculum. In the 2008 version of the software, the “Unit Conversions” unit was broken up into two smaller units; for the sake of between-year comparisons, we re-combined them.

For instance, Figure 4 suggests students that some students skipped “unit conversions” to work on “1st quadrant linear graphs”, and skipped “1-step linear equations” to work on “independent variables in linear models”, indicated by the subsequent unit being worked on by a greater proportion of students than the immediately prior unit.

[to replace last two sentences on p6 (extending to p8)]




Figure 5: The percentages of year-2 students with usage data who worked at least one problem from each unit in the Algebra I curriculum. The units are arranged in order for the standard curriculum. Students are divided between those attending schools using primarily a customized curriculum and those using primarily the standard Algebra I curriculum.

replace that preceded them in the CTAI curriculum. Most strikingly, “Lin. Eq. w/ Var’s Both Sides” was worked by a greater proportion of students in Year 2 than in Year 1, and by a greater proportion of students than any of the previous six sections. Presumably teachers and administrators wanted students to focus on that unit, perhaps because they found it to be particularly effective, because students tend to struggle with its main topic, or because its topic may figure prominently in an upcoming standardized test.

Most of the variation in unit order was driven by the rise, in year 2, of customized curricula. Figure 5 divides year-2 students into those attending schools using primarily a customized curriculum, and those attending schools using primarily a standardized curriculum.¹ Students using a standardized curriculum followed the standard sequence—more or less—while students using customized curricula did not. That said, there were some order violations in the standard group: specifically, more students worked problems from units “4-Quadrant Linear Graphs” and “Exponents” than worked the preceding sections; this suggests that some teachers used the reassignment tool to prioritize particular topics. Of

¹At least 90% of problems at “Customized” schools were from a customized curriculum, and at most 1% of problems at “Standard” schools were from a customized curriculum.

teacher reassignment

course, this may have occurred in schools with customized curricula as well—a possibility we will discuss in the next section.

Figure 6 further decomposes the year-2 results by school and state, showing a large amount of variation between states, as well as variation between schools within states. In Texas, every school used customized curricula, most of which seem to prioritize some of the same units, for instance, “Linear Patterns,” “Ind. Variables in Lin. Mod.,” and “Lin. Eq. w/ Var’s Both Side.” ~~These topics might be emphasized on Texas standardized tests.~~ On the other hand, there was also variance between schools. For instance, one school prioritized units “2 Step Lin. Eq.” and “4Quadrant Linear Graphs” while nearly eliminating “Linear Patterns.”

Between-school variation is evident in the other states, as well. In four of the five Kentucky schools, nearly every student worked on the first nine units; in the one Kentucky school that used a customized curriculum, nearly every student worked on the first 13 units, omitted the 15th (“Lin. Mod. in General Form”), and worked on the 16th and 17th (“Literal Eq.” and “Lin. Eq. w/ Var’s Both Sides”). In the remaining school, nearly every student worked on the first section, but usage decreased rapidly from there. ~~Curiously, in one Michigan school which used the standard curriculum, no students seem to have worked on the “Lin. Mod. & Ratios” section.~~

If unit order and topic scaffolding are important to CT’s mastery learning mechanism, the wide variation in students’ realized curricula would seem to pose a problem. The fact that the prescribed order was followed less in the second year of the study, when CTAI was effective, than in the first year, when it wasn’t, suggests that the standard curriculum may play a smaller role than one might otherwise imagine.

6 Mastering the Material—Or Not

The central idea behind mastery learning is that students progress through the curriculum as they master skills. In the context of CT, skills are clustered within sections, which are in turn clustered within units. Students progress from the current section to the next section after mastering all of the current section’s skills. Ideally, students would master all of the skills in all of the sections they work.

By default, the software operates by automatically moving students from section to section based on the sequence of topics defined by the curriculum they were currently enrolled in. In this software-controlled sequencing, students ideally spend the time necessary to learn the material of a section, are judged by the software to have mastered the material, and then “graduate” to the next section. Students who exhaust a section’s material without mastering its skills are promoted to the next section. ~~Students left uninterrupted follow the sequence of topics defined by their curriculum, ideally graduating from topics by showing mastery. However, teachers could modify a student’s path within the curriculum. They could “reassign” students from their current sections to other sections earlier or later in the intended sequence, including sections they had worked on previously. Finally, if the semester ends, or a student stops using CT for some other reason, while in the middle of working through a~~

Additionally, teachers are able to

However, the software will also “promote” a student to the next section if the student exhausts a section’s material without mastering its skills.

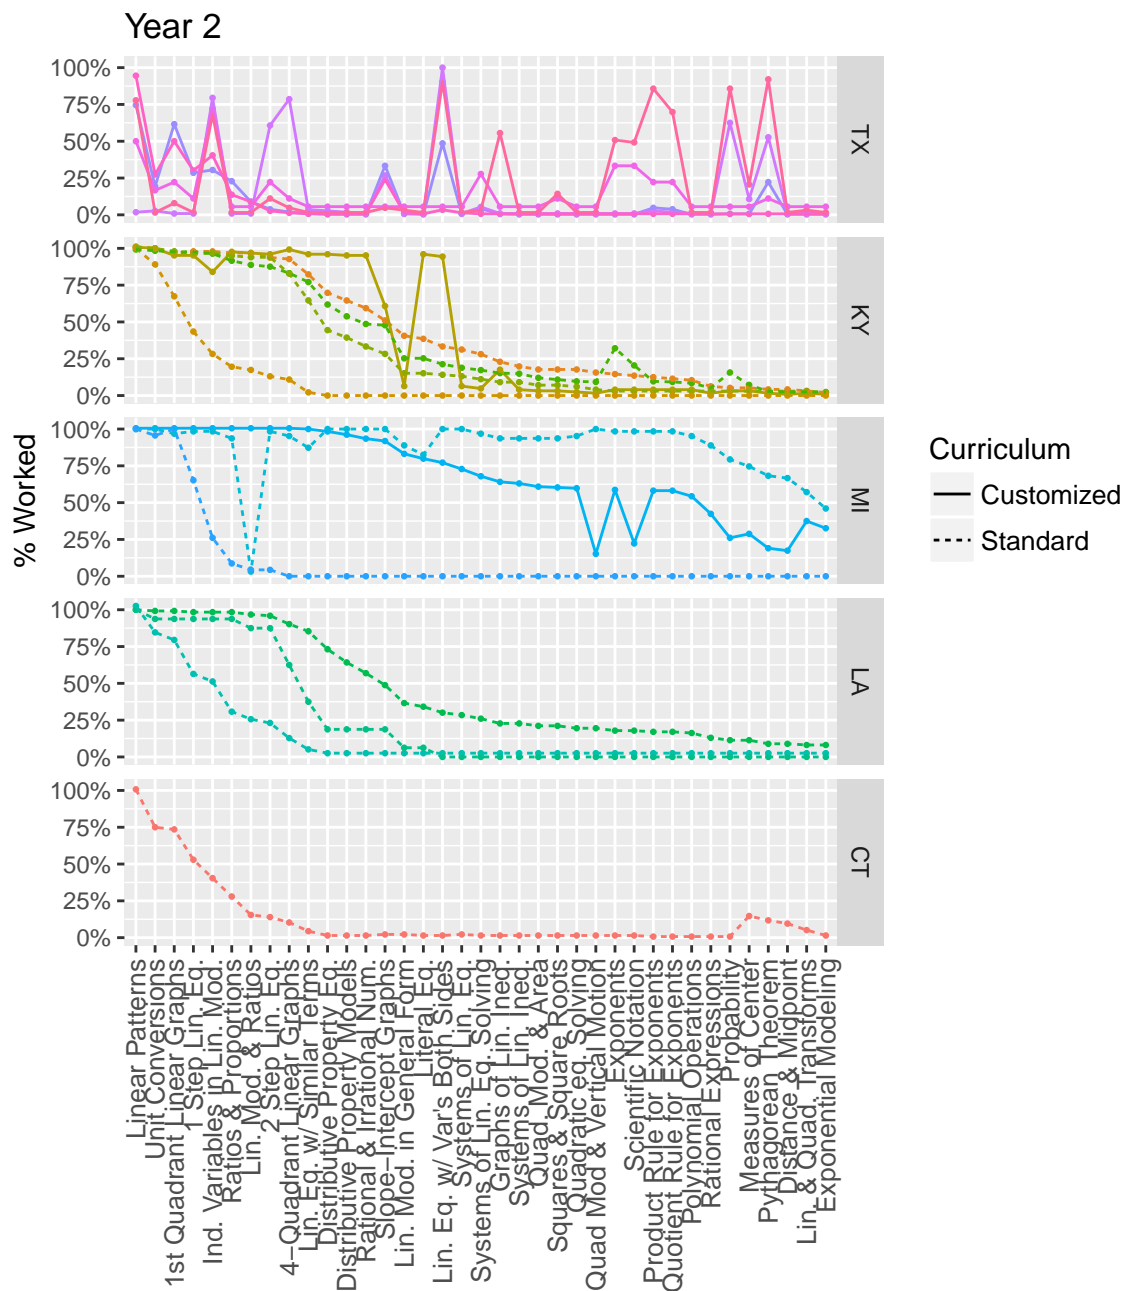


Figure 6: The percentages of year-2 students with usage data in each school who worked at least one problem from each unit in the Algebra I curriculum. The units are arranged in order for the standard curriculum. Schools are classified as either using primarily customized curricula (solid line) or using primarily the standard Algebra I curriculum (dotted).




Figure 7: The distributions of outcomes of worked sections, by state and across the entire sample, in the two study years.

	Year 1						Year 2					
	TX	KY	MI	LA	CT	Overall	TX	KY	MI	LA	CT	Overall
Final	699	625	242	129	91	1796	1252	660	236	137	114	2412
Reassigned	600	140	220	5	92	1082	1215	253	275	9	144	2042
Promoted	1619	9628	4582	1459	261	17561	1000	6296	4717	1998	450	14477
Mastered	15815	56199	31200	8648	1303	113374	4929	36676	26509	10108	1513	79874

Table 2: Percentages of worked sections that ended in each of the four possible outcomes, across states and study years.

the section is designated as "final", another situation where the student will not yet have mastered the section's skills. All in all, each CT section begins ends in one of four possible ways: mastery, promotion, reassignment, or as the student's final section.

Figure 7 and Table 2 show the proportions of worked sections in each state and study year that ended with mastery, promotion, or reassignment, or which were the student's final section. In the first year, about 85% of worked sections are mastered in every state other than Connecticut. Other than in Texas, about 13–15% of sections end in promotion. About 3 and 5% of sections in Texas and Connecticut end in reassignment, which is even rarer in the other states.

With the exception of Texas, sections in the second year tend to be completed as they were in year 1. In Texas, however, the percentage of sections ending in reassignment increased by a factor of about 5, to about 14%. The proportion of Texas sections labeled "Final" increased as well, most likely due to the decrease in usage, and thus the total number of sections encountered.

Across states, just under 1% of sections in year 1 ended in reassignment, as did roughly sections ended in reassignment at a rate of about 1% in year 1 and 2% in year 2.




Figure 8: The distributions of outcomes of worked sections, by curriculum, in the two study years. (There were no Bridge-to-Algebra sections in customized curricula in our dataset.)

2% in year 2.

6.1 Section Mastery and Curriculum

A well-designed curriculum, can, in theory, play an important role in students' attainment of mastery. Students who work on appropriate problems that build on their current set of skills should be more likely to master new skills than students working on problems above their level. What role did curriculum play in mastery during the CTAI effectiveness trial?

Figure 8 shows the proportions of worked sections that were mastered or ended in promotion, reassignment, or finality, in standard and customized versions of each CT curriculum. Mastery proportions do, indeed, depend on curriculum. Specifically, students mastered sections from more advanced curricula less frequently. Sections from the most basic curriculum, Bridge to Algebra, were mastered 92% of the time; those from Algebra I were mastered 83% of the time, and those from more advanced curricula were mastered at a rate of 75%. This is unsurprising, since more advanced curricula may be expected to be more challenging. However, it may suggest that some students studying advanced topics would fare better in more standard curricula.

Algebra I sections from customized curricula tended to end in reassignment more often than sections from the standard Algebra I curriculum (3% vs. 1, 1%). This may indicate an overall skepticism towards the Carnegie Learning standards among certain schools and teachers.

not sure I follow — this conclusion would seem to be supported by the opposite result, standard having more reassignment than customized.

Suggest moving this figure to next page.




Figure 9: Results from a set of eight multilevel logistic regressions predicting section reassignment. For each year, in the entire sample (“Overall”) and in the three states with the largest numbers of reassignments (Texas, Kentucky, and Michigan), we regressed a binary variable indicating whether a section ended in reassignment on random intercepts for school, class, student, and unit, and in the overall case, for state as well, and recorded their variance. The residual variance was set as the variance of the standard logistic distribution, $\pi/3$. These bar charts give the proportion of the total variance attributable to each random effect.

7 Digging Deeper into Section Reassignment

The proportion of worked sections in our dataset ending in reassignment was small. Nevertheless, since reassignment represents the only mechanism by which individual teachers can affect their students’ progress through the Cognitive Tutor, exploring patterns of reassignment can provide insight into how CT was used.

7.1 How Do Reassignment Patterns Vary?

Teachers alone control reassignment. Nevertheless, the factors influencing student reassignment vary at a number of levels. For instance, state and district standards may prod teachers into reassigning students to particular units. Some principals may encourage teachers to adhere to the official curriculum and avoid reassignment. Some students may be more prone to reassignment than others. Certain units in the CTAI curriculum may be harder than others, causing students to tarry and teachers to reassign. Finally, a host of other factors, at these levels and others, may spur reassignment.

To better understand the source of the variation in reassignment—what drives some, but not other, sections to end in reassignment—we fit a set of multilevel models. We fit separate

worked by students

models to data from each the three states with the highest numbers of reassignments, Texas, Kentucky, and Michigan, and in the sample as a whole, in each of the two study years, yielding a total of eight models. Each model was a logistic regression: a binary indicator for section reassignment was regressed on a random intercept for unit, as well as nested random intercepts for student, classroom, and school. Models fit to data from all six states included an additional random intercept for state.

Logistic regression can be represented in terms of an underlying latent variable Z^* : student i working section sec is reassigned when $Z_{sec,i}^* > 0$. The model for Z^* is:

$$Z_{sec,i}^* = \alpha_0 + \beta_{u[sec]} + \gamma_i + \delta_{c[i]} + \epsilon_{s[i]} + e_{sec,i}$$

Where α_0 is an overall intercept, and $\beta_{u[sec]}$, γ_i , $\delta_{c[i]}$, and $\epsilon_{s[i]}$ are random intercepts for the unit in which sec appears, for student i , for i 's classroom, and for i 's school, respectively. Again, the model fit to all six states includes an additional random intercept for state. The random intercepts are modeled as independent and normally distributed, each with its own variance. The regression error $e_{sec,i}$ is given the standard logistic distribution, with "residual" variance $\pi/3$. It is convenient to represent variance in reassignment probabilities in terms of the variance of Z^* .

Figure 9 gives the variance components estimated from these logistic regressions: variances of the random intercept terms, as a percentage of the total variance of Z^* . Overall, in both years of the study, the largest determinant of reassignment was school, accounting for 30% of the variation in year 1, and 43% in year 2. After school, state was the most important, accounting for 20 and 21% in the two years, and unit, accounting for 16 and 14%. Surprisingly, classroom and student-level factors only accounted for 7% and 3% in year 1, respectively, and 6% and 1% in year 2.* The pattern was similar in Texas, where school accounted for over half the variation in reassignment in both years, and in Michigan. In Kentucky, unit played the largest role (52%) in year 1, and classroom played the largest role in year 2 (50%). Across states and years, student-level factors never accounted for more than 5% of the variation in reassignment. Other than in Kentucky in year 2, classroom never accounted for more than 50, 12, 12, 10, 7, 6, 5, 2% of the variation.

to a lesser extent

Although teachers control reassignment, their decisions are largely dictated by broader policies, occurring at the state or school level.

determined

? appear to be

7.2 When are Students Reassigned?

The timing of reassignments can provide a window into what drives teachers' decisions to reassign students. Figure 10 shows the proportion of worked sections in each month that end in reassignment. In both years, reassignments were much more common in the second half of the school year than in the first. This may be the result of teachers learning how to use the software as the year progresses, or responding the pressure of upcoming standardized tests by accelerating students' progress and reassigning students to relevant sections.

As we've seen, reassignment was more common in year 2 than in year 1. In fact, reassignment increases fairly steadily over the entire length of the study. Through December of

* Maybe note that percentages from the state models are not directly comparable to the overall model because state accounts for some of the variance in the overall model.




Figure 10: The proportion of worked sections ending in reassignment, by month and year.

	Year 1		Year 2	
	Aug-Dec	Jan-Jun	Aug-Dec	Jan-Jun
TX	18	582	281	932
KY	9	131	52	201
MI	47	173	245	30
LA	0	5	9	0
CT	14	78	84	60
NJ	3	22	17	129
Overall	91	991	688	1352

Table 3: The number of reassignments in each state and overall in the first and second halves of the year

the first year, reassignment was rare, and from January through May of year 1, between one and two percent of sections ended in reassignment. Year 2 began where year 1 left off, with one to two percent of sections reassigned. Finally, from February through May of the second year, the rate of reassignment increased again.

Figure 10 and Table 3 decompose these trends by state. Figure 10 includes just the three states with the highest number of reassignments, Texas, Kentucky, and Michigan, and Table 3 shows data for all six states and overall. For the two study years, Figure 10 shows the proportion of all of each state's reassignments that occurred in each month of the school year. (Note that while both Figure 10 and 11 show proportions for each month, the denominators are not the same: Figure 10 gives the proportion of each month's worked sections that ended in reassignment, and Figure 11 gives the proportion of all of each state's reassignments over the course of the year, that occurred in each month.)




Figure 11: The proportion of a state's reassignments that occurred in each month of the year.

Figure 11 reveals patterns that are not apparent in Figure 10. For example, the pattern in Figure 10 seems largely driven by Texas. In contrast, the bulk of Kentucky's reassignments took place in May. A fifth of Michigan's reassignments in year 1 occurred in October; for the rest of the year, Michigan roughly followed the same rough pattern as Texas. In year 2, the vast majority of Michigan's reassignments occurred at the beginning of the year, in September and October.

7.3 Does Reassignment Depend on Classmates?

Student individuality and independence might be the most important motivating factors behind mastery learning—each student learns at his or her own pace, and struggles on a unique set of skills. Students are supposed to move through the CT curriculum independently of each other. However, reassignments give teachers the ability to override this feature, and coordinate students' progress. Teachers ^{might} can identify students who are behind their classmates and reassign them to later sections. They may also move an entire class together to a particular section or unit of interest. To what extent did these and similar considerations drive reassignment in the CTAI study?

Figure 12 addresses this question by plotting a students' classmates' statuses at the time he or she is reassigned. For each reassignment in Texas, Kentucky, and Michigan, Figure 12 plots a vertical bar colored to show the proportions of the reassigned students' classmates (represented in the usage data) who had exited the section on the same day or earlier via promotion or mastery, who had been reassigned from that section on an earlier date, on the same date, or on a later date, who mastered the section or were promoted on a later date, or who never worked the section at all. The bars are ordered according to those

3.21

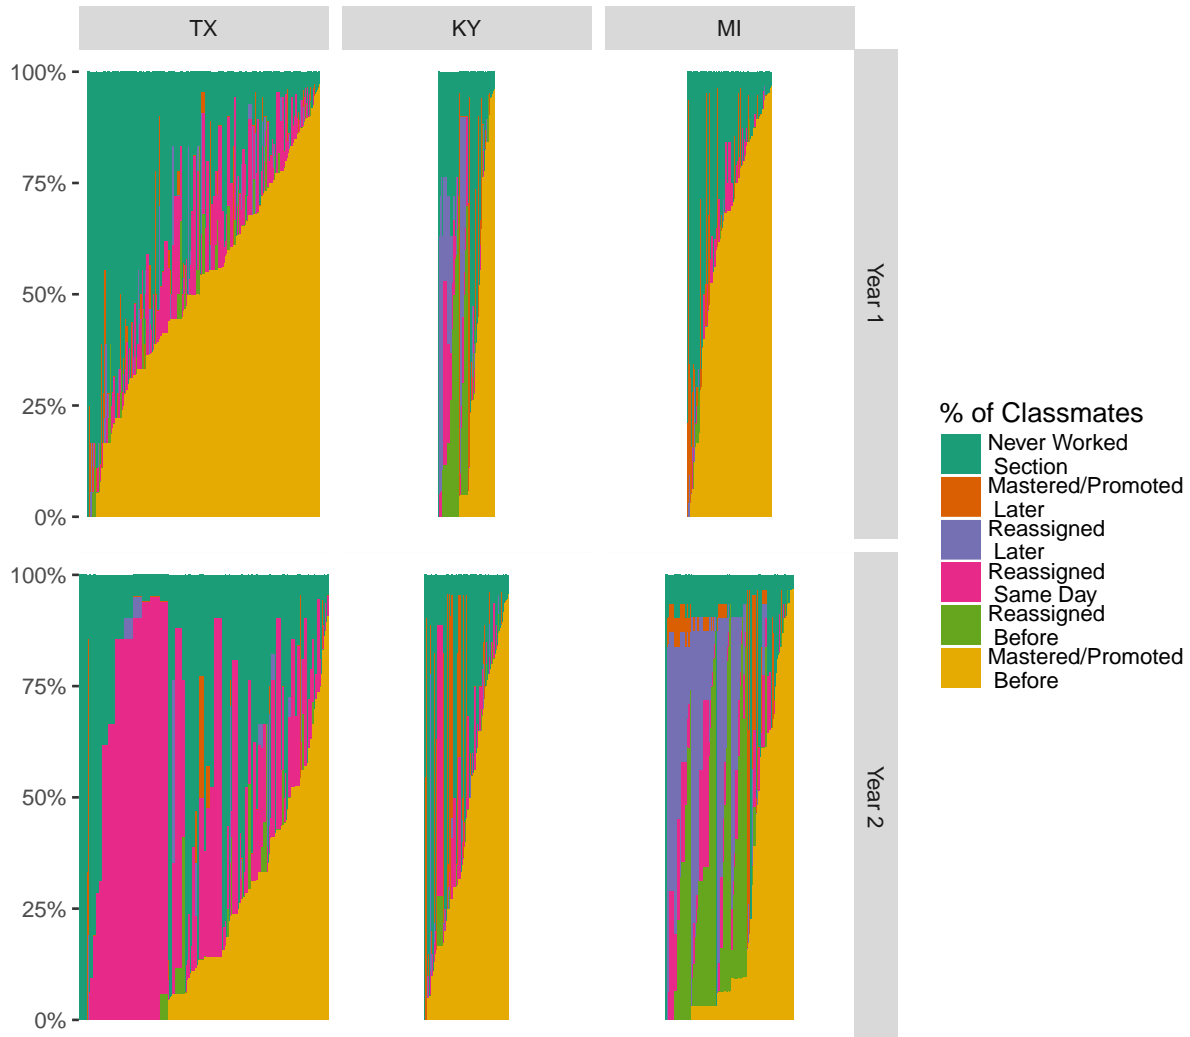


Figure 12: Distributions of classmates' statuses at each reassignment in Texas, Kentucky, and Michigan. Each reassignment that took place in these three states is represented by a bar showing the proportions of the reassigned student's classmates who had exited the section on the same day or earlier via promotion or mastery, who had been reassigned from that section on an earlier date, on the same date, or on a later date, who mastered the section or were promoted on a later date, or who never worked the section at all. The bars are ordered according to those proportions, in the order listed.

same proportions: first by the proportion of classmates who were promoted or mastered the section on the same day as the reassignment in question, or earlier, next by the proportion who were reassigned from the same section on an earlier date, next by the proportion who were reassigned from the same section on the same date, and finally by the proportion who were reassigned from the same section on a later date.

Do teachers reassign students in order to help them catch up with classmates? According to Figure 12, that might be part of the story, but isn't all of it. Across years and states, in about 40% of reassignments, at least 75% of the rest of the class had exited the section on the same date or earlier. In fact, smaller proportions of students had graduated or been promoted from the same section on the same day or earlier ^{was smaller} in year 2 than in year 1, especially in Texas and Michigan. *The figure also reveals that the*

In Texas, year 2 saw a dramatic increase in the proportions of students reassigned from the same section on the same day, suggesting that some teachers may have been moving the class together through the curriculum. In Michigan in year 2, teachers reassigned almost all students who worked certain sections. That is, ⁱⁿ 65% of the instances in which year-two Michigan students were reassigned from a section, at least 75% of their classmates were, at some point, reassigned from the same section, or never worked it. Across all three states and both years, it is exceedingly rare for students to master or be promoted from a section after someone in their class has been reassigned from the same section.

All in all, three different patterns emerge from Figure 12: teachers reassigning students who have fallen behind their classmates, teachers reassigning an entire class from the same section on the same day, and teachers reassigning almost all students who begin to work on particular sections. Each of these patterns takes place in different states and years.

7.4 Where To?

Teachers who reassign students may simply move them to the next section within the same unit. Say that a teacher believes that a particular student had already mastered the skills in one of the CTAI sections or is wheel-spinning—working problems without learning—or dislikes one of the sections in a CT unit. Still, the teacher wants the student to learn as much as possible from the ^{student} ~~student's~~ current unit. Then moving the student to the next section within the same unit ^{might} ~~would~~ make sense. Alternatively, teachers who believe that some of their students are not progressing quickly enough may reassign them out of their current units entirely, and into the next unit in the sequence—whatever that may be. Finally, a teacher who wanted his or her students to focus on a particular topic might reassign them all to the appropriate ^{unit} ~~section~~ when the time is appropriate.

Which of these patterns is most prevalent? More generally, when teachers reassign their students, ^{to} where in the curriculum do they send them?

To address these questions, we focus on the units in the standard Algebra I sequence. Figures 13 and 14 give transition plots for reassignment in the two study years. On the left of each figure, marked “From,” are the top “sending” units: the units that students were reassigned *from* at least 30 times. They are numbered according to the order they appear in the standard Algebra I curriculum. On the right, under “To,” are the top “receiving” units:

I suggest saying unit here since that is the granularity of the following discussion.

*who?
the teacher
or the
student?*

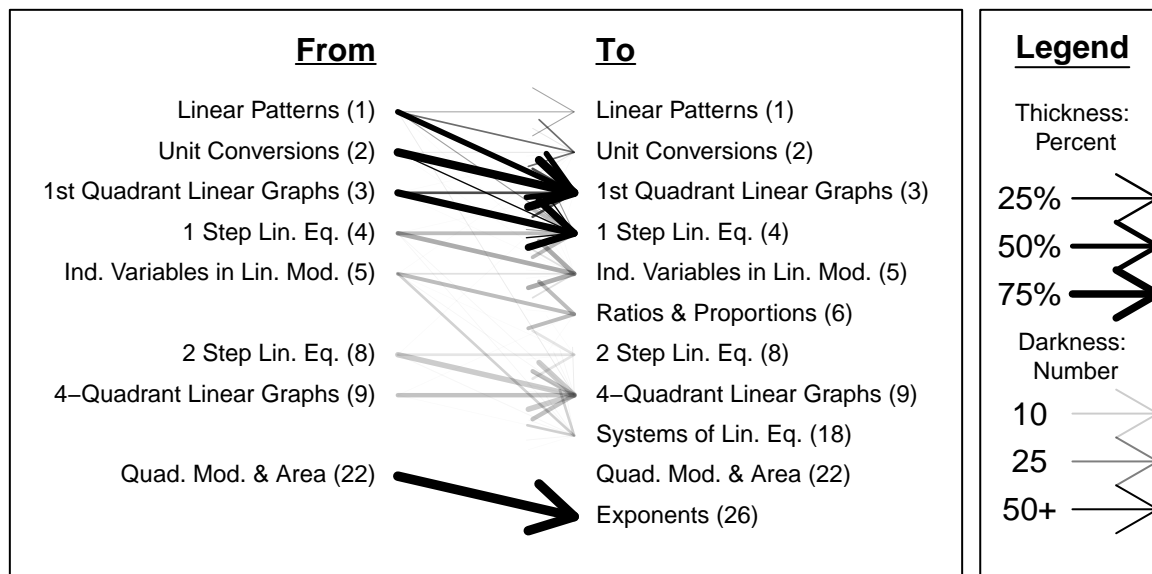


Figure 13: Reassignment transition plot for study year 1. The units on the left of the plot are the Algebra I units that ended in reassignment at least 30 times in year 1. The units on the right are those that were the destination of at least 30 reassignments. Also included on the left are the top two sending units for each of the units on the right, and also included on the right are all the units on the left, along with their top two receiving units. The units are numbered according their order in the standard Algebra I curriculum. There is an arrow from a unit on the left to a unit on the right if a student was assigned from the unit on the left to the one on the right. The thickness of the arrows is proportional to the percentage of all reassignments from the sending unit that ended in that receiving unit. The darkness of the arrows is proportional to the number of reassignments from the sending unit to the receiving unit.

these units'




Figure 14: Reassignment transition plot for study year 2. *See caption of Figure 13 for details.*

the units that students were reassigned *to* at least 30 times. For completeness, each of the top two receivers for each sending unit, and each of the top two senders for each receiving unit were also included. Finally, all of the sending units ~~were also~~ listed on the left ^{in the} receiving column. All in all, 75% of the first-year reassignments and 75% of the second-year reassignments are captured in the figures. *were also listed*

The arrows in the plot represent reassignments. An arrow from a sending unit to a receiving unit indicates reassignment from the former to the latter. The thickness of the arrows represents the proportion of reassignments originating in the sending unit whose destination was the receiving unit. The darkness of the arrows represents the number of such reassignments. For instance, in Figure 13, the arrow from “4-Quadrant Linear Graphs” on the left to “4-Quadrant Linear Graphs” on the right is fairly thick, since 44% of the reassignments begun in “4-Quadrant Linear Graphs” end in the same unit. Yet it is also fairly faint, since it only represents 12 reassignments.

Inspecting the figures shows that the most common pattern is for students to be reassigned to the next unit in the curriculum—this pattern comprises 52% of the reassignments in year 1 and 24% of those in year 2. On the other hand, it is rare for students to be reassigned within the same unit (these reassignments account for 15% and 18% in the two years). It is also rare for students to be reassigned to units earlier in the curriculum (comprising 2% and 8%).

The transition plots also reveal some interesting cases worth highlighting. In year 1 (Figure 13), students reassigned from the first unit, “Linear Patterns”, were primarily placed two units ahead, in “1st Quadrant Linear Graphs,” skipping “Unit Conversions,” perhaps suggesting a disinterest in “Unit Conversions.” on the part of the reassigning teachers. Similarly, 9 ^{of the} students (27%) ~~of the students~~ reassigned from the section “Ind. Variables in Lin. Mod.” were placed 13 units later in “Systems of Lin. Eq.” and all 62 ^{the} of students reassigned from “Quad. Mod. & Area” were placed in “Exponents.” This suggests that some teachers may have considered ~~these~~ units, “Systems of Lin. Eq.” and “Exponents” to contain particularly important material. *spell out these unit names?*

Since the total number of reassignments was higher in year 2, the corresponding plot (Figure 14) is larger and more complex. It is also more common ^{in year 2} for students to be reassigned to units other than the next unit in the sequences. This may be partly due to the proliferation of customized curricula. Two units, “4-Quadrant Linear Graphs” and “Lin. Eq. w/ Var’s Both Sides,” were common destinations from a wide variety of earlier units, suggesting ^{teacher} a strong interest in those units. The majority (76%) of the students reassigned from “Probability” were placed in another section of the same unit, perhaps suggesting problems with some of that unit’s sections; in fact, all of the students reassigned within the “Probability” unit were reassigned from one its first three sections (of seven). Finally, 44 (81%) of the reassignments from ~~sections in the~~ ^{unit} “Pythagorean Theorem” unit ended in the earlier ^{unit} “Product Rule for Exponents” sections.

In sum, the most prevalent pattern was for students to be reassigned to the following unit, suggesting that teachers ^{maybe} ~~are~~ mostly interested in ^{advancing} helping lagging students ~~progress~~. On the other hand, a number of examples of other patterns—students moving within the same

unit, or to units out of sequence—appear as well, suggesting that some teachers may be finely manipulating their students’ curricula.

8 Effects of Reassignment

The goal of CTAI is to help students learn Algebra, so the most important questions about reassignment are about its effect on learning. Although the data from this study came from a randomized trial, it was CTAI as a whole that was randomized, not individual behaviors within CTAI. Specifically, student reassignment was not randomized. Therefore, claims regarding the causal effects of reassignment on learning require strong untestable assumptions that are unlikely to be true. As in all observational studies, this includes the assumption that all confounding variables—variables that predict both reassignment and learning—have been measured well and modeled correctly. Further complicating matters, although reassignment itself is a well-defined process, in practice it can take many forms, as we have endeavored to show. There is no reason to expect the effects of reassignment to be the same regardless of whether the teacher used it to help lagging students catch up, to allocate time to important topics, or for some other reason.

All that said, observational estimates of reassignment’s average effects can be valuable, if interpreted cautiously. In the absence of evidence from randomized trials, observational studies can help guide intuition, future research, and even—when combined with other relevant information and theory—practice.

8.1 Average Effects

Figure 15 shows students gain scores—the difference between their post-test and pre-test scores—as a function of the number of times they were reassigned. (The number of reassignments was jittered—random noise was added—to avoid overplotting.) Overall, the relationship between the two variables is positive. Nevertheless, some non-linearity seems to be present, especially in year 2. Further, the distribution of the number of reassignments is right-skewed—again, especially in year 2. Care in modeling the number of reassignments, then, is especially important—observations from students reassigned an unusually large number of times can exert undue influence on a regression model and generate misleading results, particularly in the presence of non-monotonic relationships. We settled on three different strategies: first, the variable R^{bin} dichotomizes reassignment— $R^{bin} = 0$ for students who were never reassigned, and $R^{bin} = 1$ for students who were. Next, R^{cat} defines a categorical variable taking the values $R^{cat} = 0, 1, 2, 3, 4+$ for students who were reassigned 0, 1, 2, 3 or four or more times, respectively. Finally, R^{num} is the raw number of reassignments, which we include for completeness.

We used linear models to estimate the effect of reassignment, regressing post-test scores on R^{bin} , R^{cat} , or R^{num} along with fixed effects for classroom, essentially modeling reassignment as randomly assigned within classroom. Since this is unlikely to be the case (even approximately), we ran a second set of models including student level covariates: pretest




Figure 15: Students' Gain scores (post-test minus pre-test) versus the number of times the were reassigned over the course of the year (jittered), with a Loess smoother added.

	R^{bin} Effect of ≥ 1 Reassignment	Parametrization	
		R^{cat} Effect of # Reassignments:	R^{num} Effect per Reassignment:
No Covariates	-0.14 ± 0.06	1 : -0.15 ± 0.07 2 : -0.15 ± 0.1 3 : -0.07 ± 0.15 4 + : -0.12 ± 0.17	-0.04 ± 0.03
Covariate Adjusted	-0.15 ± 0.06	1 : -0.15 ± 0.07 2 : -0.16 ± 0.1 3 : -0.08 ± 0.15 4 + : -0.14 ± 0.17	-0.04 ± 0.03

Table 4: Estimates of effects of reassignment on posttests, with 95% margins of error.

scores,² race, sex, grade, special education and gifted status, English as a second language, and free and reduced-price lunch eligibility.

The results are reported in Table 4. All the effect estimates are negative, indicating that reassigning students may hurt their algebra learning. The estimated effects decrease in magnitude for three or four reassignments, but these estimates are very noisy—very few students were reassigned more than two times. The magnitudes of the effects are rather large: (?) reported an effect, in year 2, of about 0.2; the estimated effect of at least one reassignment is -15% of that. *shouldn't this be -75%?*

But what of unmeasured confounding? (?) suggest a method of estimating the sensitivity of a regression to omitted confounders based on benchmarking from observed confounders. Roughly speaking, the idea is to imagine a confounder that predicts reassignment and posttests to the same extent as one of the observed confounders, and widen the confidence interval to account for the additional uncertainty. We estimated “sensitivity intervals” allowing for the omission of a confounder similar to each of our observed covariates. *other cite?*

this paragraph seems unfinished

9 Notes

[[don't forget to include these]]

- # of students per state?
- NJ excluded from state-by-state comparisons, cuz too small. Included in “overall,” etc.
- What schools included?

²These are measured with error, and missing at a relatively high rate (15%). To account for this, we used regression calibration based on the 20 “multiple imputations” used in the original CTAI study. (?)

are these supposed to be citations of Paine et al. 2013?

Benchmark	Sensitivity	Interval
pretest	-0.15	\pm -0.04
raceBlack		\pm 0.06
raceHispanic		\pm 0.06
sexM		\pm 0.06
grade9		\pm 0.06
specspected		\pm 0.06
specgifted		\pm 0.06
specesl		\pm 0.06
fri		\pm 0.06
friMISTRUE		\pm 0.06

is table incomplete?

Table 5: Sensitivity intervals for the R^{bin} model.

- Students with no mastery data: maybe they didn't work at all, maybe they worked but we don't have their data. Treating them as unknown.